

## Neural Networks image recognition - MultiLayer Perceptron

Use both MLNN for the following problem.

1. Add random noise (see below on `size` parameter on `np.random.normal` (<https://numpy.org/doc/stable/reference/random/generated/numpy.random.normal.html>)) to the images in training and testing. *\*Make sure each image gets a different noise feature added to it. Inspect by printing out several images. Note - the `size` parameter should match the data.\**
2. Compare the `accuracy` of train and val after N epochs for MLNN with and without noise.
3. Vary the amount of noise by changing the `scale` parameter in `np.random.normal` by a factor. Use `.1`, `.5`, `1.0`, `2.0`, `4.0` for the `scale` and keep track of the `accuracy` for training and validation and plot these results.

### `np.random.normal`

#### Parameters

##### `loc`

Mean ("centre") of the distribution.

##### `scale`

Standard deviation (spread or "width") of the distribution. Must be non-negative.

##### `size`

Output shape. If the given shape is, e.g.,  $(m, n, k)$ , then  $m * n * k$  samples are drawn. If size is `None` (default), a single value is returned if `loc` and `scale` are both scalars. Otherwise, `np.broadcast(loc, scale).size` samples are drawn.

## Neural Networks - Image Recognition

```
In [37]: import keras
import numpy as np
import matplotlib.pyplot as plt
from keras.datasets import mnist
from keras.models import Sequential
from keras.optimizers import RMSprop
from keras.layers import Dense, Dropout, Flatten
from keras.layers import Conv2D, MaxPooling2D
from keras import backend
```

## Multi Layer Neural Network

Trains a simple deep NN on the MNIST dataset. Gets to 98.40% test accuracy after 20 epochs (there is a *lot* of margin for parameter tuning).

```
In [32]: # the data, shuffled and split between train and test sets
(x_train, y_train), (x_test, y_test) = mnist.load_data()

x_train = x_train.reshape(60000, 784)
x_test = x_test.reshape(10000, 784)
x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')

60000 train samples
10000 test samples
```


```
In [33]: #Define noise value & add it to the images
noise = 0.5
noise_x_train = x_train + np.random.normal(loc=0, scale=noise, size=x_train.sh
noise_x_test = x_test + np.random.normal(loc=0, scale=noise, size=x_test.shape

#Normalization of pixel values between 0 and 1 & makes them floating point num
noise_x_train = np.clip(noise_x_train, 0., 1.)
noise_x_test = np.clip(noise_x_test, 0., 1.)

#Function to compare original vs noisy images
def number_images(original, noisy, num_images=10):
    plt.figure(figsize=(20, 4))
    for i in range(num_images):
        #Original image
        ax = plt.subplot(2, num_images, i + 1)
        plt.imshow(original[i].reshape(28, 28), cmap='gray')
        plt.title("Original")
        plt.axis("Off")

        #Noisy image
        ax = plt.subplot(2, num_images, num_images + i + 1)
        plt.imshow(noisy[i].reshape(28, 28), cmap='gray')
        plt.title("Noisy")
        plt.axis("Off")
    plt.show()

#Apply the function for x_train and noise_x_train
number_images(x_train, noise_x_train)
```



In [34]: *#Question 2 NO NOISE*

```
#Define parameters for batch_size, num_classes, and epochs
batch_size = 128
num_classes = 10
epochs = 20

#Convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

#Define NN model
model = Sequential()
model.add(Dense(512, activation='relu', input_shape=(784,)))
model.add(Dropout(0.2))
model.add(Dense(512, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(10, activation='softmax'))
model.summary()

#Compile the model
model.compile(loss='categorical_crossentropy',
              optimizer="adam",
              metrics=['accuracy'])

#Fit the model on x_train + y_train & defined parameters + utilize validation
history = model.fit(x_train, y_train,
                   batch_size=batch_size,
                   epochs=epochs,
                   verbose=1,
                   validation_data=(x_test, y_test))

#Evaluation of the model + display the results
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

Model: "sequential\_8"


Layer (type)	Output Shape	Param #
dense_24 (Dense)	(None, 512)	401,920
dropout_16 (Dropout)	(None, 512)	0
dense_25 (Dense)	(None, 512)	262,656
dropout_17 (Dropout)	(None, 512)	0
dense_26 (Dense)	(None, 10)	5,130

Total params: 669,706 (2.55 MB)


Trainable params: 669,706 (2.55 MB)

Non-trainable params: 0 (0.00 B)


Epoch 1/20

469/469  2s 3ms/step - accuracy: 0.8662 - loss: 0.4456 - val\_accuracy: 0.9636 - val\_loss: 0.1058


Epoch 2/20

469/469  1s 3ms/step - accuracy: 0.9668 - loss: 0.1078 - val\_accuracy: 0.9722 - val\_loss: 0.0815


Epoch 3/20

469/469  1s 3ms/step - accuracy: 0.9776 - loss: 0.0700 - val\_accuracy: 0.9783 - val\_loss: 0.0719


Epoch 4/20

469/469  1s 3ms/step - accuracy: 0.9839 - loss: 0.0535 - val\_accuracy: 0.9782 - val\_loss: 0.0667


Epoch 5/20

469/469  1s 3ms/step - accuracy: 0.9849 - loss: 0.0469 - val\_accuracy: 0.9801 - val\_loss: 0.0676


Epoch 6/20

469/469  1s 3ms/step - accuracy: 0.9884 - loss: 0.0356 - val\_accuracy: 0.9794 - val\_loss: 0.0661


Epoch 7/20

469/469  1s 3ms/step - accuracy: 0.9877 - loss: 0.0357 - val\_accuracy: 0.9784 - val\_loss: 0.0736


Epoch 8/20

469/469  1s 3ms/step - accuracy: 0.9911 - loss: 0.0282 - val\_accuracy: 0.9812 - val\_loss: 0.0671


Epoch 9/20

469/469  1s 3ms/step - accuracy: 0.9910 - loss: 0.0260 - val\_accuracy: 0.9813 - val\_loss: 0.0692


Epoch 10/20

469/469  1s 3ms/step - accuracy: 0.9921 - loss: 0.0222 - val\_accuracy: 0.9834 - val\_loss: 0.0669


Epoch 11/20

469/469  1s 3ms/step - accuracy: 0.9910 - loss: 0.0277 - val\_accuracy: 0.9826 - val\_loss: 0.0637


Epoch 12/20

469/469  2s 3ms/step - accuracy: 0.9940 - loss: 0.0188 - val\_accuracy: 0.9830 - val\_loss: 0.0722


Epoch 13/20

469/469  2s 4ms/step - accuracy: 0.9942 - loss: 0.0170 - val\_accuracy: 0.9842 - val\_loss: 0.0669


Epoch 14/20

469/469  2s 3ms/step - accuracy: 0.9938 - loss: 0.0180 - val\_accuracy: 0.9851 - val\_loss: 0.0674


Epoch 15/20

469/469  2s 4ms/step - accuracy: 0.9942 - loss: 0.0179 - val\_accuracy: 0.9821 - val\_loss: 0.0792


Epoch 16/20

469/469  1s 3ms/step - accuracy: 0.9945 - loss: 0.0163 - val\_accuracy: 0.9838 - val\_loss: 0.0746


Epoch 17/20

469/469  1s 3ms/step - accuracy: 0.9950 - loss: 0.0151 - val\_accuracy: 0.9809 - val\_loss: 0.0834


Epoch 18/20

469/469  1s 3ms/step - accuracy: 0.9952 - loss: 0.0143 - val\_accuracy: 0.9806 - val\_loss: 0.0888

Epoch 19/20

**469/469**  **1s** 3ms/step - accuracy: 0.9950 - loss: 0.0143 -  
val\_accuracy: 0.9827 - val\_loss: 0.0824

Epoch 20/20

**469/469**  **1s** 3ms/step - accuracy: 0.9950 - loss: 0.0157 -  
val\_accuracy: 0.9837 - val\_loss: 0.0821

Test loss: 0.08285236358642578

Test accuracy: 0.9836999773979187

In [35]: `print("""After defining the parameters & not adding any noise to the images, t`

After defining the parameters & not adding any noise to the images, the test accuracy of the model turned out to be 0.9836999773979187 & the test loss for the model was 0.08285236358642578

In [39]: `###Question 2 NOISE PRESENT`

```
#Noise function that adds noise to each image & normalizes the pixel values be
def noise(images, noise_factor=0.5):
    noisy_images = images + noise_factor * np.random.normal(loc=0.0, scale=1.0)
    noisy_images = np.clip(noisy_images, 0., 1.)
    return noisy_images

#Create datasets with values that have been transformed to contain noise
noise_x_train = noise(x_train)
noise_x_test = noise(x_test)

#Define NN model with noise
noise_model = Sequential([
    Dense(512, activation='relu', input_shape=(784,)),
    Dropout(0.2),
    Dense(512, activation='relu'),
    Dropout(0.2),
    Dense(num_classes, activation='softmax')
])
model.summary()

#Compile the model
noise_model.compile(loss='categorical_crossentropy',
                    optimizer="adam",
                    metrics=['accuracy'])

#Fit the model on noise_x_train + y_train & defined parameters + utilize valid
history_noise = noise_model.fit(noise_x_train, y_train,
                                batch_size=batch_size,
                                epochs=epochs,
                                verbose=1,
                                validation_data=(noise_x_test, y_test))

#Evaluation of the model + display the results
score_noise = noise_model.evaluate(noise_x_test, y_test, verbose=0)
print('Noisy Data - Test loss:', score_noise[0])
print('Noisy Data - Test accuracy:', score_noise[1])
```

Model: "sequential\_8"

Layer (type)	Output Shape	Param #
dense_24 (Dense)	(None, 512)	401,920
dropout_16 (Dropout)	(None, 512)	0
dense_25 (Dense)	(None, 512)	262,656
dropout_17 (Dropout)	(None, 512)	0
dense_26 (Dense)	(None, 10)	5,130
dense_27 (Dense)	(None, 512)	5,632
dropout_18 (Dropout)	(None, 512)	0

dense_28 (Dense)	(None, 512)	262,656
dropout_19 (Dropout)	(None, 512)	0
dense_29 (Dense)	(None, 10)	5,130
dense_30 (Dense)	(None, 512)	5,632
dropout_20 (Dropout)	(None, 512)	0
dense_31 (Dense)	(None, 512)	262,656
dropout_21 (Dropout)	(None, 512)	0
dense_32 (Dense)	(None, 10)	5,130


**Total params:** 2,555,956 (9.75 MB)

**Trainable params:** 1,216,542 (4.64 MB)


**Non-trainable params:** 0 (0.00 B)

**Optimizer params:** 1,339,414 (5.11 MB)


Epoch 1/20

469/469  2s 3ms/step - accuracy: 0.7008 - loss: 0.8997 - val\_accuracy: 0.8781 - val\_loss: 0.3768


Epoch 2/20

469/469  1s 3ms/step - accuracy: 0.8859 - loss: 0.3457 - val\_accuracy: 0.9098 - val\_loss: 0.2859


Epoch 3/20

469/469  1s 3ms/step - accuracy: 0.9244 - loss: 0.2339 - val\_accuracy: 0.9169 - val\_loss: 0.2535


Epoch 4/20

469/469  1s 3ms/step - accuracy: 0.9470 - loss: 0.1620 - val\_accuracy: 0.9233 - val\_loss: 0.2430


Epoch 5/20

469/469  1s 3ms/step - accuracy: 0.9600 - loss: 0.1220 - val\_accuracy: 0.9267 - val\_loss: 0.2325


Epoch 6/20

469/469  1s 3ms/step - accuracy: 0.9689 - loss: 0.0950 - val\_accuracy: 0.9221 - val\_loss: 0.2775


Epoch 7/20

469/469  1s 3ms/step - accuracy: 0.9753 - loss: 0.0734 - val\_accuracy: 0.9203 - val\_loss: 0.3018


Epoch 8/20

469/469  1s 3ms/step - accuracy: 0.9759 - loss: 0.0682 - val\_accuracy: 0.9279 - val\_loss: 0.2723

Epoch 9/20

469/469  1s 3ms/step - accuracy: 0.9802 - loss: 0.0571 - val\_accuracy: 0.9269 - val\_loss: 0.2918









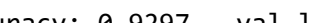
Epoch 10/20

469/469  1s 3ms/step - accuracy: 0.9820 - loss: 0.0526 - val\_accuracy: 0.9262 - val\_loss: 0.3116

Epoch 11/20

469/469  1s 3ms/step - accuracy: 0.9858 - loss: 0.0429 -



```
val_accuracy: 0.9264 - val_loss: 0.3146
Epoch 12/20
469/469  1s 3ms/step - accuracy: 0.9858 - loss: 0.0411 -
val_accuracy: 0.9270 - val_loss: 0.3188
Epoch 13/20
469/469  1s 3ms/step - accuracy: 0.9887 - loss: 0.0342 -
val_accuracy: 0.9244 - val_loss: 0.3285
Epoch 14/20
469/469  1s 3ms/step - accuracy: 0.9874 - loss: 0.0362 -
val_accuracy: 0.9266 - val_loss: 0.3447
Epoch 15/20
469/469  1s 3ms/step - accuracy: 0.9868 - loss: 0.0371 -
val_accuracy: 0.9268 - val_loss: 0.3677
Epoch 16/20
469/469  1s 3ms/step - accuracy: 0.9872 - loss: 0.0373 -
val_accuracy: 0.9265 - val_loss: 0.3671
Epoch 17/20
469/469  1s 3ms/step - accuracy: 0.9884 - loss: 0.0333 -
val_accuracy: 0.9244 - val_loss: 0.3873
Epoch 18/20
469/469  1s 3ms/step - accuracy: 0.9895 - loss: 0.0294 -
val_accuracy: 0.9241 - val_loss: 0.3819
Epoch 19/20
469/469  1s 3ms/step - accuracy: 0.9886 - loss: 0.0319 -
val_accuracy: 0.9280 - val_loss: 0.3901
Epoch 20/20
469/469  1s 3ms/step - accuracy: 0.9885 - loss: 0.0330 -
val_accuracy: 0.9297 - val_loss: 0.3531
Noisy Data - Test loss: 0.35639652609825134
Noisy Data - Test accuracy: 0.9297000169754028
```

```
In [40]: print("""By adding noise to each image (by a factor of 0.05), the test accurac
```

```
By adding noise to each image (by a factor of 0.05), the test accuracy of the
model dropped down to 0.9297000169754028, while the test loss of the model in
creased to 0.35639652609825134.
```

```
In [41]: #Define noise parameter
noise_scales = [0.1, 0.5, 1.0, 2.0, 4.0]

#Create lists to store accuracy values for both the training and validation se
accuracy_training = []
accuracy_validation = []

#Define loop that applies noise to training and testing sets for each scale in
for scale in noise_scales:
    noise_x_train = noise(x_train, scale)
    noise_x_test = noise(x_test, scale)


#Define the model
model = Sequential([
    Dense(512, activation='relu', input_shape=(784,)),
    Dropout(0.2),
    Dense(512, activation='relu'),
    Dropout(0.2),
    Dense(10, activation='softmax')
])

#Compile the model
model.compile(loss='categorical_crossentropy', optimizer="adam", metrics=[

#Train the model on noise_x_train + y_train & defined parameters + utilize
evaluation = model.fit(noise_x_train, y_train,
                        batch_size=128, epochs=20, verbose=0,
                        validation_data=(noise_x_test, y_test))

#Obtain accuracies for training and validation and append them to their re
accuracy_training.append(evaluation.history['accuracy'][-1])
accuracy_validation.append(evaluation.history['val_accuracy'][-1])

#Plot the results + analyze the results
plt.figure(figsize=(10, 5))
plt.plot(noise_scales, accuracy_training, label='Training Accuracy')
plt.plot(noise_scales, accuracy_validation, label='Validation Accuracy')
plt.title('Noise vs. Model Accuracy')
plt.xlabel('Noise Value')
plt.ylabel('Model Accuracy')
plt.legend()
plt.show()
```



```
In [45]: print("""It seems like validation accuracy is highest (>90%) when the noise va
```

It seems like validation accuracy is highest (>90%) when the noise value is  $< 0.5$ . Once the noise value is  $> 0.5$ , the validation set experiences a severe drop in accuracy.